

Project 3: CBTTC Pediatric Brain Tumor Classification

Predicting Tumor Grade from CBTTC Proteomic Profiles

This notebook is a starter project built around the Children's Brain Tumor Tissue Consortium (CBTTC) data in the Proteomic Data Commons (PDC). The baseline project uses the open-access study `PDC000180` and frames the assignment as a logistic regression classification task.

Research question: Can proteomic features from the CBTTC pediatric brain tumor study classify tumors into low-grade versus high-grade groups on test data? Does tumor grade vary by primary diagnosis? Do high-grade and low-grade tumors show different proteomic expression patterns?

Abstract

Childhood brain tumors are among the most serious forms of pediatric cancer, so improving molecular understanding of these tumors is important for diagnosis and treatment planning. This project uses the Children's Brain Tumor Tissue Consortium (CBTTC) proteomic data available through the Proteomic Data Commons to study whether tumor grade can be predicted from protein-expression measurements. The workflow downloads public study data, cleans and merges the clinical and proteomic tables, handles missing values, standardizes the numeric predictors, and trains a logistic regression classification model. Model performance is evaluated with accuracy, precision, recall, F1-score, and ROC-AUC on a held-out test set, along with supporting visualizations. The broader impact of this work is that reproducible analysis of public pediatric cancer data can help identify molecular patterns associated with more aggressive disease.

Introduction of the Data Set and Variables

The dataset used in this project comes from the **Children's Brain Tumor Tissue Consortium (CBTTC)**, a collaborative pediatric brain tumor research effort focused on identifying therapies for children with brain tumors. The data are available through the NCI Cancer Research Data Commons and the Proteomic Data Commons (PDC). For this project starter,

the selected PDC study is `PDC000180` , labeled **Pediatric Brain Cancer Pilot Study - Proteome**.

This study is connected to the pediatric brain cancer proteogenomic analysis published by Petralia et al. (2020). Tumor samples were collected through the CBTTTC and analyzed with mass-spectrometry-based proteomics, making the dataset suitable for a supervised machine-learning classification task. The modeling table in this notebook combines clinical metadata with protein abundance values, then uses tumor grade as the categorical target variable.

Sources

- CBTTTC / CRDC dataset listing: <https://datacommons.cancer.gov/explore/datasets>
- PDC study page: <https://pdc.cancer.gov/pdc/study/PDC000180>
- Petralia F. et al. (2020), *Integrated Proteogenomic Characterization across Major Histological Types of Pediatric Brain Cancer*:
<https://pmc.ncbi.nlm.nih.gov/articles/PMC8143193/>

```
In [1]: # Import the libraries used in this notebook.
import warnings
from pathlib import Path

import numpy as np
import pandas as pd
import requests
import seaborn as sns
import matplotlib.pyplot as plt

from IPython.display import display
from sklearn.decomposition import PCA
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (
    ConfusionMatrixDisplay,
    RocCurveDisplay,
    accuracy_score,
    classification_report,
    f1_score,
    precision_score,
    recall_score,
    roc_auc_score,
)
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

warnings.filterwarnings("ignore")
pd.set_option("display.max_columns", 20)
pd.set_option("display.max_rows", 200)
sns.set_theme(style="whitegrid", palette="deep")
```

```

In [2]: # Define the study ID and cache paths.
PDC_STUDY_ID = "PDC000180"
DATA_TYPE = "log2_ratio"
API_URL = "https://pdc.cancer.gov/graphql"

data_dir = Path("data")
data_dir.mkdir(exist_ok=True)
proteome_cache = data_dir / f"{PDC_STUDY_ID}_proteome.csv"
clinical_cache = data_dir / f"{PDC_STUDY_ID}_clinical.csv"

def query_pdc(query: str) -> dict:
    # Send a GraphQL query to the Proteomic Data Commons API.
    response = requests.post(API_URL, json={"query": query}, timeout=120)
    response.raise_for_status()
    payload = response.json()
    if "errors" in payload:
        raise ValueError(payload["errors"])
    return payload["data"]

def matrix_to_wide_df(matrix: list) -> pd.DataFrame:
    # Convert the PDC quantitation matrix into a gene-by-sample DataFrame.
    df = pd.DataFrame(matrix[1:], columns=matrix[0]).set_index(matrix[0][0])
    df = df.apply(pd.to_numeric, errors="coerce")
    df.columns = [str(col).split(":")[-1] for col in df.columns]
    return df.sort_index(axis=1)

def records_to_df(records) -> pd.DataFrame:
    # Convert an API record payload into a pandas DataFrame.
    if isinstance(records, list) and records and isinstance(records[0], dict):
        return pd.DataFrame(records)
    if isinstance(records, list) and records and isinstance(records[0], list):
        return pd.DataFrame(records[1:], columns=records[0])
    return pd.DataFrame(records)

def build_quant_query(study_id: str, data_type: str) -> str:
    return "\n".join(
        [
            "{",
            "    quantDataMatrix(",
            f'        pdc_study_id: "{study_id}"',
            f'        data_type: "{data_type}"',
            "        acceptDUA: true",
            "    )",
            "}"
        ]
    )

def build_clinical_query(study_id: str) -> str:
    return "\n".join(
        [

```

```

        "{",
        f'    clinicalMetadata(pdc_study_id: "{study_id}" acceptDUA: true) {{',
        "        aliquot_submitter_id",
        "        morphology",
        "        primary_diagnosis",
        "        tumor_grade",
        "        tumor_stage",
        "    }",
        "}",
    ]
)

def map_grade_group(value):
    # Map detailed grade labels into low-grade and high-grade groups.
    if pd.isna(value):
        return np.nan

    text = str(value).strip().lower()
    high_tokens = ["high", "grade iii", "grade iv", "grade 3", "grade 4", "g3", "g4"]
    low_tokens = ["low", "grade ii", "grade i", "grade 2", "grade 1", "g1", "g2"]

    if any(token in text for token in high_tokens):
        return "High Grade"
    if any(token in text for token in low_tokens):
        return "Low Grade"
    return np.nan

```

```

In [3]: # Download the CBTC proteomic matrix and clinical metadata, or use cached files if
if proteome_cache.exists() and clinical_cache.exists():
    proteome_wide = pd.read_csv(proteome_cache, index_col=0)
    clinical_df = pd.read_csv(clinical_cache)
else:
    proteome_response = query_pdc(build_quant_query(PDC_STUDY_ID, DATA_TYPE))
    clinical_response = query_pdc(build_clinical_query(PDC_STUDY_ID))

    proteome_wide = matrix_to_wide_df(proteome_response["quantDataMatrix"])
    clinical_df = records_to_df(clinical_response["clinicalMetadata"])

    proteome_wide.to_csv(proteome_cache)
    clinical_df.to_csv(clinical_cache, index=False)

print("Proteome matrix shape (proteins x samples):", proteome_wide.shape)
print("Clinical metadata shape:", clinical_df.shape)

display(clinical_df.head())
display(proteome_wide.iloc[:5, :5])

```

Proteome matrix shape (proteins x samples): (7154, 230)
 Clinical metadata shape: (293, 5)

	aliquot_submitter_id	morphology	primary_diagnosis	tumor_grade	tumor_stage
0	7316-324	9505/1	Ganglioglioma, NOS	G1	Unknown
1	7316-445	9380/3	Glioma, malignant	High Grade	Unknown
2	7316-445	9380/3	Glioma, malignant	High Grade	Unknown
3	7316-329	9400/3	Glioma, NOS	G1	Unknown
4	7316-206	9400/3	Glioma, NOS	G1	Unknown

7316-101 7316-109 7316-111 7316-114 7316-117

Gene/Aliquot					
A1BG	-1.7408	-0.7403	-0.3692	0.9485	-0.0447
A2M	-1.0443	-0.7829	-0.3893	0.1037	0.0394
AAAS	NaN	NaN	NaN	-0.2482	-0.6455
AACS	0.1790	0.4741	NaN	-0.0699	NaN
AADAT	NaN	NaN	NaN	NaN	NaN

```
In [4]: # Display the dataset shape and the data types of the clinical variables.
clinical_df = clinical_df.drop_duplicates(subset="aliquot_submitter_id").copy()

dataset_overview = pd.DataFrame(
    {
        "Table": ["Clinical metadata", "Proteome matrix"],
        "Rows": [clinical_df.shape[0], proteome_wide.shape[0]],
        "Columns": [clinical_df.shape[1], proteome_wide.shape[1]],
    }
)

clinical_dtypes = pd.DataFrame(
    {
        "Variable": clinical_df.columns,
        "Data Type": clinical_df.dtypes.astype(str).values,
    }
)

display(dataset_overview)
display(clinical_dtypes)
print("Proteomic data types:", proteome_wide.dtypes.unique())
```

	Table	Rows	Columns
0	Clinical metadata	227	5
1	Proteome matrix	7154	230

	Variable	Data Type
0	aliquot_submitter_id	object
1	morphology	object
2	primary_diagnosis	object
3	tumor_grade	object
4	tumor_stage	object

Proteomic data types: [dtype('float64')]

Data Cleaning and Preprocessing

The steps below align the clinical and proteomic data by sample ID, remove duplicates, inspect missing values, convert the categorical target into a binary outcome, and prepare a modeling table for classification.

```
In [5]: # Convert the proteome matrix so that each row is one tumor sample, then merge with
proteome_by_sample = proteome_wide.T.copy()
proteome_by_sample.index.name = "aliquot_submitter_id"

merged_df = clinical_df.merge(
    proteome_by_sample,
    left_on="aliquot_submitter_id",
    right_index=True,
    how="inner",
)

merged_df["grade_group"] = merged_df["tumor_grade"].apply(map_grade_group)
model_df = merged_df.dropna(subset=["grade_group"]).copy()

clinical_columns = [
    "aliquot_submitter_id",
    "morphology",
    "primary_diagnosis",
    "tumor_grade",
    "tumor_stage",
    "grade_group",
]
protein_columns = [col for col in model_df.columns if col not in clinical_columns]

missing_summary = model_df.isna().sum().sort_values(ascending=False)
protein_missing_rate = model_df[protein_columns].isna().mean().sort_values(ascending=True)
kept_proteins = protein_missing_rate[protein_missing_rate <= 0.20].index.tolist()

encoded_metadata = pd.get_dummies(
    model_df[["morphology", "primary_diagnosis", "tumor_stage"]],
    dummy_na=True,
    drop_first=False,
)
```

```

print("Merged table shape:", model_df.shape)
print("Protein features before filtering:", len(protein_columns))
print("Protein features kept for modeling:", len(kept_proteins))
print("\nTarget distribution:")
print(model_df["grade_group"].value_counts())

display(missing_summary.head(10).to_frame("missing_values"))
display(encoded_metadata.head())
display(model_df[clinical_columns].head())

```

Merged table shape: (218, 7160)
 Protein features before filtering: 7154
 Protein features kept for modeling: 4958

Target distribution:
 grade_group
 Low Grade 159
 High Grade 59
 Name: count, dtype: int64

missing_values	
PRR4	210
DHRS2	209
CILP	208
RCAN3	199
PNMT	199
IMPG1	198
MMP10	198
RSPH14	194
IZUMO1	193
PIWIL2	192

	morphology_9350/1	morphology_9380/3	morphology_9391/3	morphology_9400/3	morphology_9400/3
0	False	False	False	False	False
1	False	True	False	False	False
3	False	False	False	True	True
4	False	False	False	True	True
5	False	False	False	True	True



	aliquot_submitter_id	morphology	primary_diagnosis	tumor_grade	tumor_stage	grade_g
0	7316-324	9505/1	Ganglioglioma, NOS	G1	Unknown	Low G
1	7316-445	9380/3	Glioma, malignant	High Grade	Unknown	High G
3	7316-329	9400/3	Glioma, NOS	G1	Unknown	Low G
4	7316-206	9400/3	Glioma, NOS	G1	Unknown	Low G
5	7316-207	9400/3	Glioma, NOS	G1	Unknown	Low G

Visualization

These plots summarize the target balance, the most common diagnoses, reduced-dimension proteomic structure, and the most variable proteins in the CBTTTC modeling subset.


```
In [6]: # Plot class balance and the most frequent diagnoses.
viz_df = model_df.copy()
viz_df["primary_diagnosis"] = viz_df["primary_diagnosis"].fillna("Unknown")
top_diagnoses = viz_df["primary_diagnosis"].value_counts().head(10).index

fig, axes = plt.subplots(1, 2, figsize=(16, 5))

sns.countplot(
    data=model_df,
    x="grade_group",
    order=model_df["grade_group"].value_counts().index,
    ax=axes[0],
)
axes[0].set_title("Tumor Grade Distribution")
axes[0].set_xlabel("Grade Group")
axes[0].set_ylabel("Number of Samples")

sns.countplot(
    data=viz_df[viz_df["primary_diagnosis"].isin(top_diagnoses)],
    y="primary_diagnosis",
    order=viz_df["primary_diagnosis"].value_counts().head(10).index,
    ax=axes[1],
)
axes[1].set_title("Top Diagnoses in the Modeling Subset")
axes[1].set_xlabel("Number of Samples")
axes[1].set_ylabel("Primary Diagnosis")

plt.tight_layout()
plt.show()
```

```
In [7]: # Visualize the proteomic data with PCA and a heatmap of the most variable proteins
viz_features = model_df[kept_proteins].copy()
viz_imputed = pd.DataFrame(
    SimpleImputer(strategy="median").fit_transform(viz_features),
    columns=kept_proteins,
    index=model_df.index,
)
viz_scaled = StandardScaler().fit_transform(viz_imputed)

pca = PCA(n_components=2, random_state=42)
components = pca.fit_transform(viz_scaled)
pca_df = pd.DataFrame(components, columns=["PC1", "PC2"], index=model_df.index)
pca_df["grade_group"] = model_df["grade_group"].values

top_variable_proteins = viz_imputed.var().sort_values(ascending=False).head(15).index
heatmap_df = viz_imputed[top_variable_proteins].copy()
heatmap_df["grade_group"] = model_df["grade_group"].values
heatmap_df = heatmap_df.sort_values("grade_group").drop(columns="grade_group")
heatmap_df = (heatmap_df - heatmap_df.mean()) / heatmap_df.std(ddof=0)


fig, axes = plt.subplots(1, 2, figsize=(16, 6))

sns.scatterplot(
    data=pca_df,
    x="PC1",
    y="PC2",
    hue="grade_group",
    s=90,
    alpha=0.9,
    ax=axes[0],
)
axes[0].set_title("PCA of Proteomic Features")
axes[0].set_xlabel(f"PC1 ({pca.explained_variance_ratio_[0]:.1%})")
axes[0].set_ylabel(f"PC2 ({pca.explained_variance_ratio_[1]:.1%})")

sns.heatmap(
    heatmap_df.T,
    cmap="vlag",
    center=0,
    xticklabels=False,
    cbar_kws={"label": "Z-score"},
    ax=axes[1],
)
axes[1].set_title("Top 15 Most Variable Proteins")
```

```
axes[1].set_xlabel("Samples")
axes[1].set_ylabel("Protein")

plt.tight_layout()
plt.show()
```



Research Question #2, Does tumor grade vary by primary diagnosis?

```
In [8]: # read file, clinical
df = pd.read_csv("PDC000180_clinical(in).csv")
df.head()
```

```
Out[8]:
```

	aliquot_submitter_id	morphology	primary_diagnosis	tumor_grade	tumor_stage
0	7316-324	9505/1	Ganglioglioma, NOS	G1	Unknown
1	7316-445	9380/3	Glioma, malignant	High Grade	Unknown
2	7316-445	9380/3	Glioma, malignant	High Grade	Unknown
3	7316-329	9400/3	Glioma, NOS	G1	Unknown
4	7316-206	9400/3	Glioma, NOS	G1	Unknown

```
In [9]: # show info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 293 entries, 0 to 292
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   aliquot_submitter_id  293 non-null   object
1   morphology            293 non-null   object
2   primary_diagnosis     293 non-null   object
3   tumor_grade           293 non-null   object
4   tumor_stage           293 non-null   object
dtypes: object(5)
memory usage: 11.6+ KB
```

```
In [10]: # read file: proteome
df = pd.read_csv("PDC000180_proteome(in).csv")
df.head()
```

```
Out[10]:
```

	Gene/Aliquot	7316-101	7316-109	7316-111	7316-114	7316-117	7316-119	7316-120	7316-121.1	7316-121.2
0	A1BG	-1.7408	-0.7403	-0.3692	0.9485	-0.0447	1.0001	-0.7616	1.5691	1.6100
1	A2M	-1.0443	-0.7829	-0.3893	0.1037	0.0394	-0.4928	-0.7111	0.2410	0.3280
2	AAAS	NaN	NaN	NaN	-0.2482	-0.6455	NaN	-1.1472	0.0826	NaN
3	AACS	0.1790	0.4741	NaN	-0.0699	NaN	-0.1300	0.5614	NaN	-0.0930
4	AADAT	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

5 rows × 231 columns




```
In [11]: # show info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7154 entries, 0 to 7153
Columns: 231 entries, Gene/Aliquot to 7316-99
dtypes: float64(230), object(1)
memory usage: 12.6+ MB
```

```
In [12]: # make a countplot to answer research question: Does tumor grade vary by primary di
plt.figure(figsize = (10,5))
sns.countplot(data=clinical_df, x='primary_diagnosis', hue = 'tumor_grade')
plt.xticks(rotation=45)
plt.title("Tumor Grade by Primary Diagnosis")
plt.xlabel("Primary Diagnosis")
plt.ylabel("Count")

plt.show()
```



Explanation: This graph shows the distribution of tumor grades and how they vary across different primary diagnoses. As shown, the x-axis is the primary diagnosis like Glioma, NOS. The legend is the tumor grade itself, such as G1, G2, G4, and high grade/low grade. The y-axis, count, is the number of patients. The graphs shows that tumor types like Glioma, NOS, have both low-grade and high-grade tumors. While others are more concentrated in a certain category. This could mean that tumor type may be associated with how severe the tumor is.

Research question #3, Do high-grade and low-grade tumors show different proteomic expression patterns?


```
In [13]: # selects protein columns
protein_cols= [col for col in model_df.columns
               if col not in ['aliquot_submitter_id', 'morphology', 'primary_d',
                             'tumor_grade', 'tumor_stage', 'grade_group']]

proteins = protein_cols[:5]

subset = model_df[proteins + ['grade_group']].melt(
    id_vars = 'grade_group',
    value_vars = proteins,
    var_name = 'Protein',
    value_name = 'Expression'
)
```

```
#create boxplot
plt.figure(figsize=(10,6))
sns.boxplot(data=subset, x='Protein', y='Expression', hue = 'grade_group')
plt.title("Protein Expression Differences: High vs Low Grade")
plt.ylabel("Expression Level")
plt.xlabel("Proteins")
plt.xticks(rotation=45)

plt.show()
```



Explantion: This boxplot compares the protein expression levels between low-grade and high-grade tumors. The x-axis is the proteins, A1BG, A2M etc, while the y-axis is the expression level. The legend helps indicate which tumors are low grade and which tumors are high grade. The boxplot shows that high-grade tumors tend to have high expression levels than low-grade tumors. There are a few outliers, however overall this graph could suggest that certain proteins vary by tumor grade. This helps support the idea that proteomic features can help distinguish between tumor severity.

Input Preparation and Normalization/Standardization

This is a classification problem, so the target variable does not require outlier treatment. The code below creates `X` and `y`, splits the data into training and test sets using stratification, and standardizes the numeric input features within a machine-learning pipeline.

```
In [14]: # Create the feature matrix and binary target for logistic regression.
X = model_df[kept_proteins].copy()
y = model_df["grade_group"].map({"Low Grade": 0, "High Grade": 1})

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.25,
    random_state=42,
    stratify=y,
)

candidate_k = sorted({min(k, len(kept_proteins)) for k in [10, 20, 40, 80, 120]})
candidate_k = [k for k in candidate_k if k > 0]

print("Training set shape:", X_train.shape)
print("Test set shape:", X_test.shape)
print("Candidate feature counts:", candidate_k)
```

Training set shape: (163, 4958)

Test set shape: (55, 4958)

Candidate feature counts: [10, 20, 40, 80, 120]

Model Building and Training

The baseline model is a logistic regression pipeline with median imputation, standardization, univariate feature selection, and hyperparameter tuning by cross-validation.

```
In [15]: # Train and tune the logistic regression model.
logistic_pipeline = Pipeline(
    steps=[
        ("imputer", SimpleImputer(strategy="median")),
        ("scaler", StandardScaler()),
        ("selector", SelectKBest(score_func=f_classif, k=candidate_k[0])),
        (
            "model",
            LogisticRegression(
                max_iter=5000,
                class_weight="balanced",
                solver="liblinear",
            ),
        ),
    ],
)

param_grid = {
    "selector__k": candidate_k,
    "model__C": [0.01, 0.1, 1, 10, 100],
}

grid_search = GridSearchCV(
    estimator=logistic_pipeline,
```

```

    param_grid=param_grid,
    scoring="roc_auc",
    cv=5,
    n_jobs=-1,
)

grid_search.fit(X_train, y_train)

best_model = grid_search.best_estimator_
selected_mask = best_model.named_steps["selector"].get_support()
selected_features = X_train.columns[selected_mask]
feature_scores = pd.Series(
    best_model.named_steps["selector"].scores_,
    index=X_train.columns,
).sort_values(ascending=False)

print("Best cross-validated ROC-AUC:", round(grid_search.best_score_, 4))
print("Best parameters:", grid_search.best_params_)
display(feature_scores.head(15).to_frame("ANOVA_F_score"))

```

Best cross-validated ROC-AUC: 0.9359

Best parameters: {'model__C': 0.01, 'selector__k': 10}

ANOVA_F_score	
CTPS1	206.425641
RCC1	179.779026
XPO5	173.864427
HNRNPAB	171.098520
FEN1	166.510425
MCM2	157.048886
MCM6	154.180949
SET	152.646580
PLEC	151.364614
DNAJC9	145.520424
MCM4	144.059558
NXF1	144.011560
NASP	142.582558
RCC2	142.244373
KPNA2	140.240585

Performance Evaluation on Test Data

This section reports the required classification metrics on the held-out test set and visualizes model performance using a confusion matrix and ROC curve.

```
In [16]: # Evaluate the tuned model on test data.
y_pred = best_model.predict(X_test)
y_prob = best_model.predict_proba(X_test)[:, 1]

metrics_df = pd.DataFrame(
    {
        "Metric": ["Accuracy", "Precision", "Recall", "F1-score", "AUC"],
        "Value": [
            accuracy_score(y_test, y_pred),
            precision_score(y_test, y_pred, zero_division=0),
            recall_score(y_test, y_pred, zero_division=0),
            f1_score(y_test, y_pred, zero_division=0),
            roc_auc_score(y_test, y_prob),
        ],
    }
)

display(metrics_df.round(4))
print(
    classification_report(
        y_test,
        y_pred,
        target_names=["Low Grade", "High Grade"],
        zero_division=0,
    )
)

fig, axes = plt.subplots(1, 2, figsize=(14, 5))


ConfusionMatrixDisplay.from_predictions(
    y_test,
    y_pred,
    display_labels=["Low Grade", "High Grade"],
    cmap="Blues",
    ax=axes[0],
)
axes[0].set_title("Confusion Matrix")

RocCurveDisplay.from_predictions(y_test, y_prob, ax=axes[1])
axes[1].set_title("ROC Curve")

plt.tight_layout()
plt.show()
```


	Metric	Value
0	Accuracy	0.8727
1	Precision	0.7500
2	Recall	0.8000
3	F1-score	0.7742
4	AUC	0.8950

	precision	recall	f1-score	support
Low Grade	0.92	0.90	0.91	40
High Grade	0.75	0.80	0.77	15
accuracy			0.87	55
macro avg	0.84	0.85	0.84	55
weighted avg	0.88	0.87	0.87	55



Conclusion

The logistic regression model used standardized proteomic features from the CBTTTC pediatric brain tumor study to classify tumors into low-grade and high-grade groups. Based on the held-out test results, the model showed how strongly protein expression patterns can separate tumor severity in this pediatric dataset. The selected proteins can be discussed as potentially informative biomarkers, while the main limitations are sample size, class imbalance, and study-specific generalizability. The second research question asked: “Does tumor grade vary by primary diagnosis?” The countplot displays multiple bars, representing the different tumor grades (x-axis). While the y-axis shows the number of patients (count). The graph concludes that tumor grade does vary by primary diagnosis. Certain diagnoses, like Glioma, NOS, have a much higher number of lower-grade tumors like G1 or G4. However other tumor grades, like Ependymoma, NOS, show higher counts in more complex grades, like G2. This infers that the type of tumor diagnosis is correlated to how severe or complex

the tumor is. Overall, concluding that tumor grade does vary by the primary diagnosis, as not all of the bars on the countplot are even, indicating differences in diagnosis and complexities. The third research question asked: "Do high-grade and low-grade tumors show different proteomic expressions?" To best answer this question, a boxplot was created to show the distributions of tumor grades and their protein expressions. The boxplot shows that high-grade and low-grade tumors have different expression levels. Many of the proteins involving high-grade tumors, had higher expression levels compared to low-grade tumors. The boxplot showed several different medians as well which shows that different protein expression varies depending on the severity of the tumor. This means that proteomic expression is related to tumor grade, while it can also help explain biological differences between less aggressive and more aggressive tumors. In addition to the data findings, mindfulness is always important to take into consideration. Patterns were recognized and interpretations occurred regarding tumor grade, diagnosis, and protein expression, however it's important to remember that this is all information on real patients. Therefore, it's crucial to respect the privacy of each patient, as this dataset contains sensitive information. In addition to that, it's always important to be careful with conclusions that are being made, as this data set contains missing values which could affect the outcomes of this research. Future steps within this research could include expanding this analysis by using a larger dataset and examining additional proteins to better understand tumor behavior. While building predictive models using this data could help identify tumor grade or severity at an early stage, improving diagnosis and treatment planning. In conclusion, this study demonstrates that proteomic characteristics can be used to classify pediatric brain tumors as low-grade or high-grade. This model performed well, demonstrating that protein expression contains useful information about tumor severity. Moreover, tumor grade depends on the initial diagnosis, and there were clear differences in proteomic expression across tumor types. Despite some shortcomings, such as sample size and class imbalance, these results suggest that proteomics could be a promising avenue for further research.

In []: